

# Final Project Submission

- Student name: Kevin Spring
- Student pace: Flex
- Scheduled project review date/time: August 8, 2022 13:30 - 14:15 (CDT)
- Instructor name: Morgan Jones
- Blog post URL: <https://medium.com/@kevinjspring/predicting-home-sale-prices-using-linear-regression-4ebf079a48e8>

## Summary

- Our client wants to be able to predict sales price, identify where to market in King County, WA, and how customers can improve their home to increase sale price.
- Ordinary least squares linear regression was used to create three models.
- The three models were compared using  $R^2$ , Prediction Intervals (PI), and Root Mean Squared Error (RMSE).
- Model 2 (M2) is the best model as it has the best predictive capabilities, R-squared of 0.88, low RMSE and PI, though the error is not normally distributed.

## Actionable Recommendations

1. Improving the condition of a house by one-unit will increase the sale price by about 6%.
2. Adding an additional full bathroom would increase the sale price of a house by about 3.9%.
3. Marketing should be focused throughout King County, WA except in Zip codes 98002, 98003, 98023, 98032, 98042, and 98198 as the value of the homes sold in these Zip codes are well under the rest of King County, WA.
4. Market real estate services toward owners of waterfront properties as these sell for 59% more than homes not waterfront.

## Table of Contents

- [Business Problem](#)
- [Stakeholders](#)
- [Background](#)
- [Data](#)
- [Data Cleanup](#)
- [Exploratory Data Analysis](#)
- [Model Specification](#)
- [Modeling and Regression Results](#)

- [Conclusions](#)
  - [Model Analysis and Comparisons](#)
  - [Prediction Application](#)
- [Recommendations](#)
- [References](#)

## Business Problem

Our client is a residential real estate broker in King County, WA interested in finding a solution for their customers. Many of their customers come to them needing to sell their home but are unsure of the market value. The client wants us to design and implement a model where they can take in the features of a seller's home and determine which price to begin listing discussion.

## Stakeholders

- President of Bon Jovi Real Estate Advisors
- Bon Jovi real estate agents

## Background

Our client wants us to predict a continuous value, sales price, from features of the house their customer gives them in the form of continuous and categorical data. Regression analysis is a statistical process to estimate the relationship between a dependent variable (response) and a continuous independent variables (predictors).

In this analysis I will use ordinary least squares (OLS) linear regression to assess the relationship between features of homes and sale price. OLS fits a linear model on data by minimizing the sum of the squared difference between the observed dependent variable and the predicted response ( $\hat{y}$ )

To calculate  $\hat{y}$ ,

$$\hat{y} = \hat{\beta}_0 + \sum_{i=1}^n x_n \hat{\beta}_n$$

where  $n$  is the number of predictors,  $\beta_0$  is the intercept,  $\hat{x}_n$  is the  $n^{th}$  predictor, and  $\hat{y}$  are the predicted value associated with the dependent variables.

The linear equation that is returned can be used to predict the response value using new data.

To perform OLS linear regression the data needs to be clean with no missing values and categorical data needs to be coded correctly. The assumptions of OLS linear regression are then checked and models are built. These models are compared using, coefficient of determination,

prediction intervals, and Root Mean Squared Error to compare and determine which model is the most suited for our client.

```
In [1]: # Import Libraries
        from datetime import date

        ## Data analysis
        import pandas as pd
        import numpy as np

        ## Statistical analysis
        from scipy import stats
        from scipy.stats import norm
        import statsmodels.api as sm
        import statsmodels.formula.api as smf
        from statsmodels.formula.api import ols

        ## Model Validation
        from sklearn.linear_model import LinearRegression
        from sklearn.model_selection import KFold
        from sklearn.model_selection import cross_val_score

        ## Visualization
        import matplotlib.pyplot as plt
        import seaborn as sns
        %matplotlib inline

        # import data
        df = pd.read_csv('data/kc_house_data.csv')
```

# Data

## Description

The data is a collection of single family homes in the King County, WA area sold between May 2014 and May 2015 (1). The data contains 21 variables and 21,597 records. This data will be suitable to create a model to predict sale price for homes within the parameters of this dataset.

```
In [2]: df.shape
```

```
Out[2]: (21597, 21)
```

```
In [3]: #inspect data
        df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21597 entries, 0 to 21596
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                     21597 non-null  int64
1   date                   21597 non-null  object
2   price                  21597 non-null  float64
3   bedrooms               21597 non-null  int64
4   bathrooms              21597 non-null  float64
5   sqft_living            21597 non-null  int64
6   sqft_lot               21597 non-null  int64
7   floors                 21597 non-null  float64
8   waterfront             19221 non-null  object
9   view                   21534 non-null  object
10  condition              21597 non-null  object
11  grade                  21597 non-null  object
12  sqft_above             21597 non-null  int64
13  sqft_basement          21597 non-null  object
14  yr_built               21597 non-null  int64
15  yr_renovated           17755 non-null  float64
16  zipcode                21597 non-null  int64
17  lat                    21597 non-null  float64
18  long                   21597 non-null  float64
19  sqft_living15          21597 non-null  int64
20  sqft_lot15             21597 non-null  int64
dtypes: float64(6), int64(9), object(6)
memory usage: 3.5+ MB

```

## Table 1 Variable Names and Descriptions for King County Data Set

See the [King County Assessor Website](#) for further explanation of each condition code

Variable	Data Type	Description
id	catagorical	Unique identifier for a house
date	continuous	Date house was sold
price	continuous	Sale price (prediction target)
bedrooms	discrete	Number of bedrooms
bathrooms	discrete	Number of bathrooms
sqft_living	continuous	Square footage of living space in the home
sqft_lot	continuous	Square footage of the lot
floors -	discrete	Number of floors (levels) in house
waterfront	ordinal	Whether the house is on a waterfront
view	ordinal	Quality of view from house
condition	ordinal	How good the overall condition of the house is. Related to maintenance of house

Variable	Data Type	Description
grade	ordinal	Overall grade of the house. Related to the construction and design of the house
sqft_above	continuous	Square footage of house apart from basement
sqft_basement	continuous	Square footage of the basement
yr_built	catagorical	Year when house was built
yr_renovated	catagorical	Year when house was renovated
zipcode	catagorical	ZIP Code used by the United States Postal Service
lat	catagorical	Latitude coordinate
long	catagorical	Longitude coordinate
sqft_living15	continuous	The square footage of interior housing living space for the nearest 15 neighbors
sqft_lot15	continuous	The square footage of the land lots of the nearest 15 neighbors

## Location of King County, WA home sales

```
In [4]: ## Map of home sales between May 2014 and May 2015
# code adapted from
# Ahmed Qassim,
# https://towardsdatascience.com/easy-steps-to-plot-geographic-data-on-a-map-python-11

# Define bounding box
BBox = ((df.long.min(), df.long.max(),
        df.lat.min(), df.lat.max() ))

# Make scatterplot
fig, ax = plt.subplots(figsize = (13,12))
ax.scatter(df.long, df.lat, c = np.log(df.price), alpha=.075,
          s=20, edgecolors='none',
          cmap= plt.cm.get_cmap('jet_r'))

# Plot paramaters
ax.set_title('King County, WA home sales between May 2014 - May 2015') # title
# Remove x, y ticks and labels
ax.tick_params(axis='both', which='both',
              bottom=False, top=False, left=False, right=False,
              labelbottom=False, labeltop=False, labelleft=False, labelright=False)

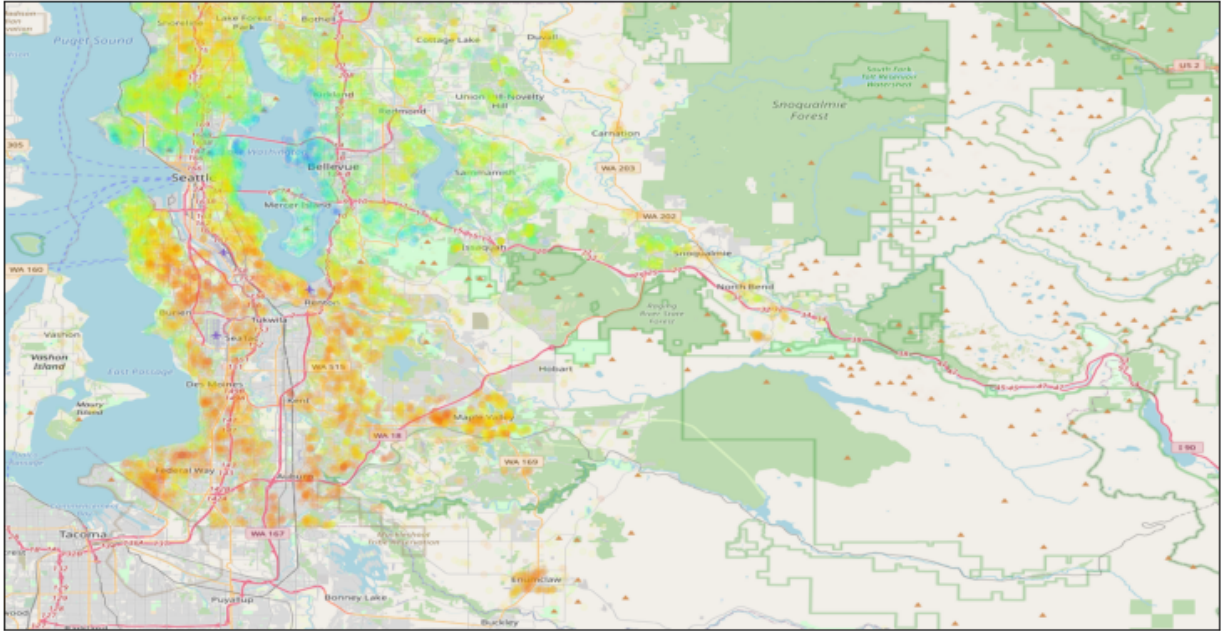
# Set x and y-axis limits to bounding box
ax.set_xlim(BBox[0],BBox[1])
ax.set_ylim(BBox[2],BBox[3])

# Set area map
ruh_m=plt.imread('img/King_County_map.png')
ax.imshow(ruh_m, zorder=0, extent = BBox, aspect= 'equal') #

# plt.savefig('img/KC_home_sale_map.png', dpi=600) # save the image
```

```
Out[4]: <matplotlib.image.AxesImage at 0x200c7b8a820>
```

King County, WA home sales between May 2014 - May 2015



## Data Limitations

- Data is only from 2014 to 2015. Models to predict future sales price would need to be updated with newer data.
- Some data might be missing, such as for-sale-by-owner or owner-financed sales.
- Ordinal data might be highly variable based on examiner's subjective experience.
- As the map shows, home sales are a mix of urban and rural houses, but much more homes are clustered together. The models may not be able to accurately predict rural house prices because of the lack of data for rural homes.

## Data Cleanup

### Identify and remove duplicated records

```
In [5]: # Any duplicated homes?
duplicates_len = len(df[df.duplicated(subset=['id'],
                                     keep=False)].sort_values(by='id'))

print(f"Results:\nThere are {duplicates_len} duplicated records.")
df[df.duplicated(subset=['id'], keep=False)].sort_values(by='id').head(4)
```

Results:  
There are 353 duplicated records.

Out[5]:

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront
<b>2495</b>	1000102	4/22/2015	300000.0	6	3.0	2400	9373	2.0	NC
<b>2494</b>	1000102	9/16/2014	280000.0	6	3.0	2400	9373	2.0	NaN
<b>16800</b>	7200179	10/16/2014	150000.0	2	1.0	840	12750	1.0	NC
<b>16801</b>	7200179	4/24/2015	175000.0	2	1.0	840	12750	1.0	NC

4 rows × 21 columns

## Duplicate home ID discussion

The duplicated records based on ID are from the same homes that sold within the same year. These homes have the same attributes except for sale `date`. These may be homes that were flipped or sold quickly after an initial sale. I will keep these records as I am interested in predicting a home's sale price and these give more data for the true value of a house.

## Remove Unnecessary variables

The following variables will be deleted from this analysis as they are unnecessary to my analysis.

- `id` - This is an unique identifier for each home. Too unique.
- `date` - This is the sale date and time will not be analyzed due to the single year of the data.
- `lat` - This is the latitude of the home sold. Will use Zip code for location.
- `long` - same reasoning as `lat`

```
In [6]: # delete unnecessary columns
df.drop(['id', 'date', 'lat', 'long'], axis=1, inplace=True)
```

## Identify Missing data

```
In [7]: # How many columns have NaN?
print(df.isna().sum())
```

```
price            0
bedrooms         0
bathrooms        0
sqft_living      0
sqft_lot         0
floors           0
waterfront      2376
view             63
condition        0
grade           0
sqft_above       0
sqft_basement    0
yr_built         0
yr_renovated     3842
zipcode          0
sqft_living15    0
sqft_lot15       0
dtype: int64
```

```
In [8]: # Any placeholders?
# Look for top occurring values
print('King County, WA \n Home Sales Dataframe\n')
for col in df.columns:
    print(col, '\n', df[col].value_counts(normalize = True).head(10), '\n')
```



King County, WA  
Home Sales Dataframe

price  
450000.0 0.007964  
350000.0 0.007964  
550000.0 0.007362  
500000.0 0.007038  
425000.0 0.006945  
325000.0 0.006853  
400000.0 0.006714  
375000.0 0.006390  
300000.0 0.006158  
525000.0 0.006066  
Name: price, dtype: float64

bedrooms  
3 0.454878  
4 0.318655  
2 0.127796  
5 0.074131  
6 0.012594  
1 0.009075  
7 0.001760  
8 0.000602  
9 0.000278  
10 0.000139  
Name: bedrooms, dtype: float64

bathrooms  
2.50 0.248970  
1.00 0.178312  
1.75 0.141131  
2.25 0.094782  
2.00 0.089364  
1.50 0.066907  
2.75 0.054869  
3.00 0.034866  
3.50 0.033847  
3.25 0.027272  
Name: bathrooms, dtype: float64

sqft\_living  
1300 0.006390  
1400 0.006251  
1440 0.006158  
1800 0.005973  
1660 0.005973  
1010 0.005973  
1820 0.005927  
1480 0.005788  
1720 0.005788  
1540 0.005742  
Name: sqft\_living, dtype: float64

sqft\_lot  
5000 0.016576  
6000 0.013428  
4000 0.011622  
7200 0.010187

```
4800    0.005510
7500    0.005510
4500    0.005279
8400    0.005140
9600    0.005047
3600    0.004769
Name: sqft_lot, dtype: float64
```

```
floors
1.0    0.494189
2.0    0.381303
1.5    0.088438
3.0    0.028291
2.5    0.007455
3.5    0.000324
Name: floors, dtype: float64
```

```
waterfront
NO      0.992404
YES     0.007596
Name: waterfront, dtype: float64
```

```
view
NONE      0.901923
AVERAGE   0.044441
GOOD       0.023591
FAIR       0.015325
EXCELLENT  0.014721
Name: view, dtype: float64
```

```
condition
Average    0.649164
Good       0.262861
Very Good  0.078761
Fair       0.007871
Poor       0.001343
Name: condition, dtype: float64
```

```
grade
7 Average    0.415521
8 Good       0.280826
9 Better     0.121082
6 Low Average 0.094365
10 Very Good 0.052507
11 Excellent 0.018475
5 Fair       0.011205
12 Luxury    0.004121
4 Low        0.001250
13 Mansion   0.000602
Name: grade, dtype: float64
```

```
sqft_above
1300    0.009816
1010    0.009724
1200    0.009538
1220    0.008890
1140    0.008520
1400    0.008334
1060    0.008242
1180    0.008196
```

```
1340    0.008149
1250    0.008057
Name: sqft_above, dtype: float64
```

```
sqft_basement
0.0    0.593879
?      0.021021
600.0   0.010048
500.0   0.009677
700.0   0.009631
800.0   0.009307
400.0   0.008520
1000.0  0.006853
900.0   0.006575
300.0   0.006575
Name: sqft_basement, dtype: float64
```

```
yr_built
2014    0.025883
2006    0.020975
2005    0.020836
2004    0.020049
2003    0.019447
2007    0.019308
1977    0.019308
1978    0.017919
1968    0.017641
2008    0.016993
Name: yr_built, dtype: float64
```

```
yr_renovated
0.0    0.958096
2014.0  0.004112
2013.0  0.001746
2003.0  0.001746
2007.0  0.001690
2000.0  0.001633
2005.0  0.001633
2004.0  0.001239
1990.0  0.001239
2009.0  0.001183
Name: yr_renovated, dtype: float64
```

```
zipcode
98103    0.027874
98038    0.027272
98115    0.026994
98052    0.026578
98117    0.025605
98042    0.025328
98034    0.025235
98118    0.023475
98023    0.023105
98006    0.023059
Name: zipcode, dtype: float64
```

```
sqft_living15
1540    0.009122
1440    0.009029
1560    0.008890
```

```
1500    0.008334
1460    0.007825
1580    0.007733
1610    0.007686
1720    0.007686
1800    0.007686
1620    0.007594
Name: sqft_living15, dtype: float64
```

```
sqft_lot15
5000    0.019771
4000    0.016484
6000    0.013335
7200    0.009724
4800    0.006714
7500    0.006575
8400    0.005371
3600    0.005140
4500    0.005140
5100    0.005047
Name: sqft_lot15, dtype: float64
```

## Missing value results

- NaN
  - waterfront
    - Binary categorical variable ( YES or NO )
    - replace NaN with mode of NO as most likely these properties are not waterfront
  - view
    - Ordinal categorical variable
    - replace NaN with NONE
  - yr\_renovated
    - Will be converted to a countable numerical variable
    - 0 is the most common value with over 95% of values.
    - Replace NaN with 0 value
- Placeholder
  - yr\_renovated has 0 for missing or unknown values.
  - sqft\_basement has ? for missing or unknown values.

```
In [9]: # replacing waterfront NaN with 'NO'
df['waterfront'].fillna('NO', inplace=True)

# replace yr_renovated NaN with 'Unknown'
df['yr_renovated'].fillna(0, inplace=True)

# replace `NaN` with `NONE` for column `view`
df['view'].fillna('NONE', inplace=True)
```

```
In [10]: # Confirm no more NaN values
print(df.isna().sum())
```

```

price          0
bedrooms       0
bathrooms      0
sqft_living    0
sqft_lot       0
floors         0
waterfront     0
view           0
condition      0
grade          0
sqft_above     0
sqft_basement  0
yr_built       0
yr_renovated   0
zipcode        0
sqft_living15  0
sqft_lot15     0
dtype: int64

```

**Table 2: Coding ordinal, binary, and count data**

variable	Data Type	Plan
condition	ordinal	Recode to dictionary. { 'Poor': 0, 'Fair': 1, 'Average': 2, 'Good': 3, 'Very Good': 4 }
grade	ordinal	Delete the descriptor, keep the number, and convert it to <code>int</code> datatype. Example: 7 Average becomes 7
basement	binary	If there is a basement (sq.ft > 0) the value will be set to 1 . No basement (sq.ft = 0) set to 0 . ? makes up about 2% of values and the current value of 0 makes up almost 60%. Replace ? with the mode of 0 .
view	ordinal	Recode to dictionary. { 'NONE': 0, 'FAIR': 1, 'AVERAGE': 2, 'GOOD': 3, 'EXCELLENT': 4 }
waterfront	binary	Recode to dictionary. { 'NO': 0, 'YES': 1 } .
home_age	discrete	Create variable from <code>yr_built</code> . Subtract current year from <code>yr_built</code> . Drop <code>yr_built</code>
yr_since_reno	discrete	Create variable from <code>yr_renovated</code> . Subtract current year from <code>yr_renovated</code> . 0 is the most common value with over 95% of values. If never renovated then subtract from <code>yr_built</code> . Drop <code>yr_renovated</code> .

```

In [11]: # -----#
# Encoding ordinal, binary, and count variables
# Code condition to ordinal data

```

```

# Map condition variable to dictionary
condition_dict = {'Poor': 0, 'Fair': 1, 'Average': 2, # Map
                  'Good': 3, 'Very Good': 4}
df['condition'] = df['condition'].map(condition_dict) # Use map to
                                                    # code values

# Code Grade to ordinal data
# Strip out by spaces and keep the first string, which is the value
df['grade'] = df['grade'].apply(lambda x: x.split(' ', 1)[0]).astype(int)

# Code sqft_basement to binary data
# sqft_basement has '?' as a placeholder. Set this to 0.
df['sqft_basement'].replace('?', 0, inplace=True)
# change to numerical type
df['sqft_basement'] = df['sqft_basement'].astype(float)
# With a basement then code as 1
df['sqft_basement'].loc[df['sqft_basement'] > 0] = 1
# rename column
df.rename(columns={'sqft_basement': 'basement'}, inplace=True)

# Code view to ordinal data
# Map ordinal variable to dictionary
view_dict = {'NONE': 0,
             'FAIR': 1,
             'AVERAGE': 2,
             'GOOD': 3,
             'EXCELLENT': 4} # map
df['view'] = df['view'].map(view_dict) # Recode

# Recode waterfront to binary data
# Map binary variable to dictionary
waterfront_dict = {'NO': 0, 'YES': 1} # map
df['waterfront'] = df['waterfront'].map(waterfront_dict) # Recode

# Recode home_age to discrete data
# Calculate home age
current_year = date.today().year # assign current year
df['home_age'] = current_year - df['yr_built'] # Calculate year since built
df.drop('yr_built', axis=1, inplace=True) # drop old column

# Recode yr_since_reno to discrete data
# subtraction function
def sub(a, b):
    return a - b
# Calculate years since last renovation
df['yr_since_reno'] = df.apply(
    lambda row : sub(current_year, row['yr_renovated']) # subtract
    if row['yr_renovated'] > 0 # if the property has been renovated
    else row['home_age'], axis = 1) # else the property has not been renovated
df.drop('yr_renovated', axis=1, inplace=True)

```

D:\ProgramData\Anaconda3\lib\site-packages\pandas\core\indexing.py:1732: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
self._setitem_single_block(indexer, value, name)
```

# Outliers

```
In [12]: # Data description to identify outliers
df.describe().apply(lambda s: s.apply(lambda x: format(x, 'g')))
```

```
Out[12]:
```

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition
count	21597	21597	21597	21597	21597	21597	21597	21597	21597
mean	540297	3.3732	2.11583	2080.32	15099.4	1.4941	0.0067602	0.233181	2.4
std	367368	0.926299	0.768984	918.106	41412.6	0.539683	0.0819439	0.764673	0.6
min	78000	1	0.5	370	520	1	0	0	1
25%	322000	3	1.75	1430	5040	1	0	0	1
50%	450000	3	2.25	1910	7618	1.5	0	0	1
75%	645000	4	2.5	2550	10685	2	0	0	1
max	7.7e+06	33	8	13540	1.65136e+06	3.5	1	4	4

There is an outlier that may be due to a data entry mistake. One house has 33 bedrooms. I was expecting it to be a mansion but it has an average grade (7), 1.75 bathrooms, and only 1,620 square feet of living space. I think this house had a miskey and the number of bedrooms should be 3.

```
In [13]: df[df['bedrooms'] > 30]
```

```
Out[13]:
```

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition	grade
15856	640000.0	33	1.75	1620	6000	1.0	0	0	4	7

```
In [14]: df.at[15856, 'bedrooms'] = 3
df.loc[[15856]]
```

```
Out[14]:
```

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition	grade
15856	640000.0	3	1.75	1620	6000	1.0	0	0	4	7

## Exploratory Data Analysis

### Cleaned Data Description

```
In [15]: df.describe().apply(lambda s: s.apply(lambda x: format(x, 'g')))
```

Out[15]:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condo
<b>count</b>	21597	21597	21597	21597	21597	21597	21597	21597	21597
<b>mean</b>	540297	3.37181	2.11583	2080.32	15099.4	1.4941	0.0067602	0.233181	2.4
<b>std</b>	367368	0.904096	0.768984	918.106	41412.6	0.539683	0.0819439	0.764673	0.6
<b>min</b>	78000	1	0.5	370	520	1	0	0	0
<b>25%</b>	322000	3	1.75	1430	5040	1	0	0	0
<b>50%</b>	450000	3	2.25	1910	7618	1.5	0	0	0
<b>75%</b>	645000	4	2.5	2550	10685	2	0	0	0
<b>max</b>	7.7e+06	11	8	13540	1.65136e+06	3.5	1	4	4

## About the data

The median house sold in King County, WA between 2014 to 2015 was for \$450,000. The median house sold was 1910 square feet, 3 bedroom, 2.25 bathrooms, and 47 years old. The home sale price range was \$78,000 to \$7,700,000.

## Variable scatter matrix

A scatter matrix will return both histograms and scatterplots lined up for each variable. The diagonal represents the histograms of that variable to visualize the distribution and the rest are scatterplots to visualize the relationships between variables.

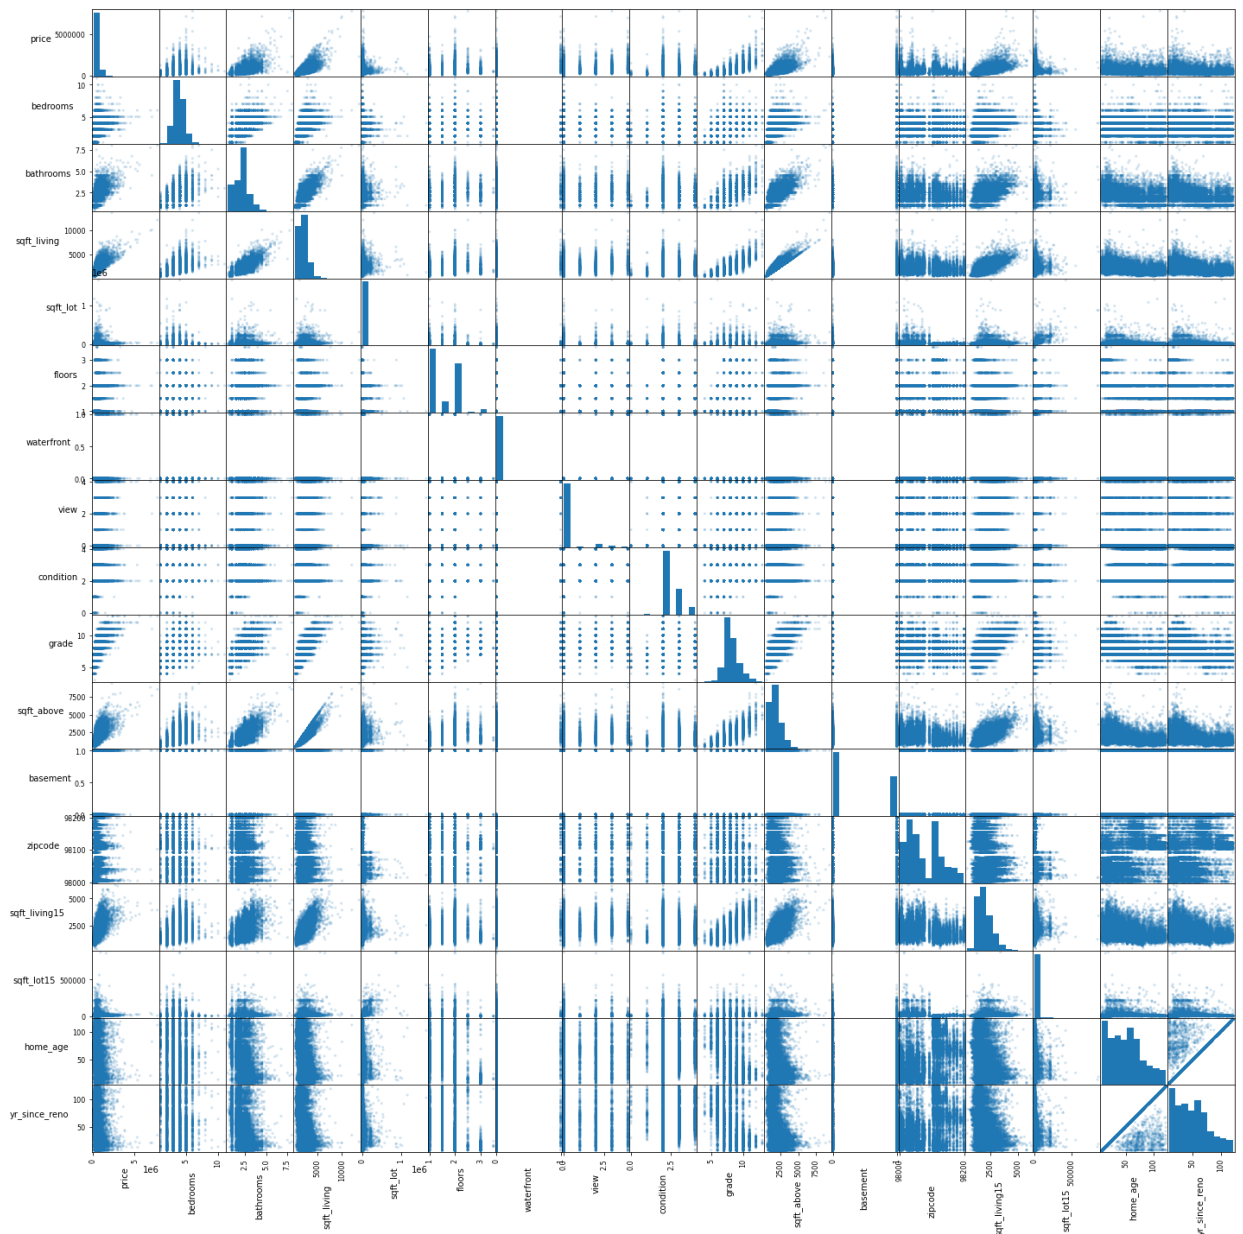
```
In [16]: # Visualize the data using scatter plot and histogram

import warnings
warnings.filterwarnings('ignore') # Ignore warnings

# Create scatter matrix
axes = pd.plotting.scatter_matrix(df, alpha = 0.2, figsize = [20, 20])
for ax in axes.flatten():
    ax.xaxis.label.set_rotation(90)
    ax.yaxis.label.set_rotation(0)
    ax.yaxis.label.set_ha('right')

plt.tight_layout()
plt.gcf().subplots_adjust(wspace=0, hspace=0)
plt.show()
```





## Scatter Matrix Results

### Histogram

The diagonal plots are the histogram and indicate that most of the variables are right-skewed, including the dependent variable, price.

### Scatterplot

Looking at the first row, the variables with the strongest positive correlation with price are for the number of bathrooms, grade, and square footage of the living space in the house.

## Correlation Matrix Heatmap

A correlation heatmap calculates the Pearson correlation between variables. A Pearson correlation measures the relationship between two variables. A value of 1 means a complete

positive correlation, a value of 0 means no correlation, and -1 means a negative correlation exists. This heatmap has a dark color for negatively correlated variables and a light color for positively correlated variables.


```
In [17]: # Make heatmap
# Code adapted from Flatiron Data Science

# compute the correlation matrix
corr = df.corr()

# Set up figure and axes
fig, ax = plt.subplots(figsize=(15, 10))

# Plot a heatmap of the correlation matrix, with both
# numbers and colors indicating the correlations
sns.heatmap(
    # Specifies the data to be plotted
    data=corr,
    # The mask means we only show half the values,
    # instead of showing duplicates. It's optional.
    mask=np.triu(np.ones_like(corr, dtype=bool)),
    # Specifies that we should use the existing axes
    ax=ax,
    # Specifies that we want labels, not just colors
    annot=True,
    # Customizes colorbar appearance
    cbar_kws={"label": "Correlation", "orientation": "horizontal", "pad": .2, "extend"
)

# Customize the plot appearance
ax.set_title("Heatmap of Correlation Between Attributes (Including Target)");
```



## Dummy variables for Zip code

Dummy variables or One-hot-encoding is a way to use categorical variables with regression analysis. The variable `zipcode` is a categorical variable and must be converted to a numerical data. Each Zip code will become its own variable and be either a 'no' (0) or 'yes' (1).

```
In [18]: # Convert zip code variable to Dummy variables
df_clean = df.copy()
cat_col = ['zipcode']

# Label columns as category
df[cat_col] = df[cat_col].astype('category')
ohe_df = pd.get_dummies(df[cat_col], drop_first=True)

# merge ohe_df with df_clean and drop old zip_code column
df_clean = pd.concat([df_clean, ohe_df], axis=1)
df_clean.drop(cat_col, axis=1, inplace=True)
```

## Distirubtion of price



```
In [19]: def hist_plot(data, Y):
    '''
    Histogram plot function
    Input:
        data: pandas dataframe
        Y: column of variable for the histogram
    Output:
        Histogram plot
        Skewness and kurtosis value
    Citation:
        Atanu Dan
        https://medium.com/@atanudan/kurtosis-skew-function-in-pandas-aa63d72e20c
    '''

    y = data[Y]
    # Plot code
    fig, ax = plt.subplots(1,2, figsize=(15,4))
    sns.distplot(y, fit=norm, bins=30, kde=False, ax=ax[0]);
    ax[0].title.set_text(f'Histogram of {Y}')
    ax[0].set(xlabel=f'{Y}', ylabel='frequency')
    res = stats.probplot(df_clean['price'], plot=ax[1])

    ## Skewness and Kurtosis
    print(f'EDA of {Y} variable')
    print(f'Skewness: {y.skew()}')
    print(f'Kurtosis: {y.kurt()}')
```

```
In [20]: #histogram of price
hist_plot(df_clean, 'price')
```

EDA of price variable  
Skewness: 4.023364652271239  
Kurtosis: 34.54135857673376



## Interpretation

The data in the variable `price` is highly right-skewed and does not follow a normal distribution as shown in the histogram and QQ-plot. This may result in a high level of **heteroskedasticity** because there are many orders of magnitude between the lowest and highest sale price. Heteroskedasticity results when variance is not equal across the range of the dependent variable. This may cause higher variance in high sale price houses in contrast to low sale price houses. In other words, the variance is unequal as it is changing proportionally with the variable.

## Model Specification

In this section I will specify which variable to include in the model to predict sales price.

Model	Description
M1	Use variables highly correlated ( $>0.6$ ) with sale price
M2	A backward stepwise regression to choose variables. This procedure will result in almost all independent variables, excluding some the dummy variables for <code>zipcode</code> , being chosen.
M3	M1 with interaction effects

## Table 3: OLS Model Assumptions (2)

Assumption	Description
1	The regression model is linear in the coefficients and error term
2	There is a random sampling of observations
3	Error term has a population mean of zero
4	There is no multi-collinearity (or perfect collinearity)
5	The error term has a constant variance (no heteroskedasticity)
6	The error term is normally distributed. This allows statistical hypothesis testing to be done.

Assumption 2 is met with the collection of the data. The data may lack private sales by owner but the majority of house sales occur through real estate agents. Assumptions 1, 3, 4, 5 and 6 will be checked with plots.

## Specifying Model 1 (M1)

I will specify M1 variables using the independent variables with a Pearson's correlation with `price` of 0.6 or greater.

```
In [21]: # Features correlated
# Source doe from Flatiron
features = []
correlations = []
for idx, correlation in corr['price'].T.iteritems():
    if correlation >= .6 and idx != 'price':
        features.append(idx)
        correlations.append(correlation)
corr_with_price = pd.DataFrame({'Correlations':correlations, 'Features': features})

print('Table 4: Independent Variables Highly Correlated (>0.6) With Price')
display(corr_with_price)
```

Table 4: Independent Variables Highly Correlated (>0.6) With Price

	Correlations	Features
0	0.701917	sqft_living
1	0.667951	grade
2	0.605368	sqft_above

## Specifying Model 2 (M2)

M2 uses an automated stepwise backward elimination feature selection strategy. All the variables are fed into the model and fitted. The independent variable with highest p-value is removed if the p-value is greater than 0.05. A p-value greater than 0.05 indicates that variable's effect is not statistically significant. This is repeated until all the p-values of the predictor variables are less than 0.05, meaning they are significant.

```
In [22]: #Backward Elimination function
def backward_elimination(df, y):
    '''
    Backward Elimination
    Feed all the possible features to the model at first. We check
    the performance of the model and then iteratively remove the worst
    performing features one by one till the overall performance of the
    model comes in acceptable range. The performance metric used here to
    evaluate feature performance is p-value. If the pvalue is above 0.05
    then we remove the feature, else we keep it.

    Input:
        df = Pandas dataframe of your data
```

```

        y = dependent variable
Output:
        list of independent variables in the model

Citation:
        Abhini Shetye
        https://towardsdatascience.com/feature-selection-with-pandas-e3690ad8504b
...
X = df.copy()
X.drop(y.columns, axis=1, inplace=True)
cols = list(X.columns) # get all of the column names
pmax = 1
while (len(cols)>0): # while there are entries in the cols list
    p= [] # initialize p-value list
    X_1 = X[cols] # create new dataframe
    X_1 = sm.add_constant(X_1) # add constant for y-intercept
    model = sm.OLS(y,X_1).fit() # OLS regression model on the data
    p = pd.Series(model.pvalues.values[1:],index = cols) # save the p-values
    pmax = max(p) # assign maximum p-valuse
    feature_with_p_max = p.idxmax() # get that p-value's index
    if(pmax > 0.05): # check if the max p-value is greater than 5
        cols.remove(feature_with_p_max) # if it is then remove it from the cols li
    else:
        break # otherwise all p-values are less than 0.05
return cols # return the features

```

## Specifying Model 3 (M3)

Interaction effects occur when the effect of one predictor variable depends on the value of another variable. For example, condition of a home may be dependent on the age of the home. There may be a dependency between the view and whether the house is on waterfront property. M3 builds on M1 by adding interaction effects to the main effects. It also removes `sqft_above` as a predictor variable due to high correlation between `sqft_above` and `sqft_living`.

The graphs below are interaction plots. This plot displays the fitted values of the dependent variable ( `price` ) on the y-axis and one of the predictor variables on the x-axis. The lines between the points represents the other predictor variable. If the lines remain parallel to each other then there is not an interaction between these variables. If the lines cross that indicates there could be an interaction.

## Interaction Effect Plot Helper Function

```

In [23]: def interaction_analysis(data, var1, var2):
...
        Produces an interaction plot with programically determined
        title.
        Input:
            data: Pandas dataframe
            var1: column name of first variable
            var2: column name of second variable
        Output:
            interaction plot with title using var1 and var2

```

```
'''
# Import Library
from statsmodels.graphics.factorplots import interaction_plot

# Make interaction plot
fig = interaction_plot(data[var1], df_clean[var2], data['price'])
plt.title(f'Interaction plot\n {var1} and {var2}') # set title
```

## Interaction Plots

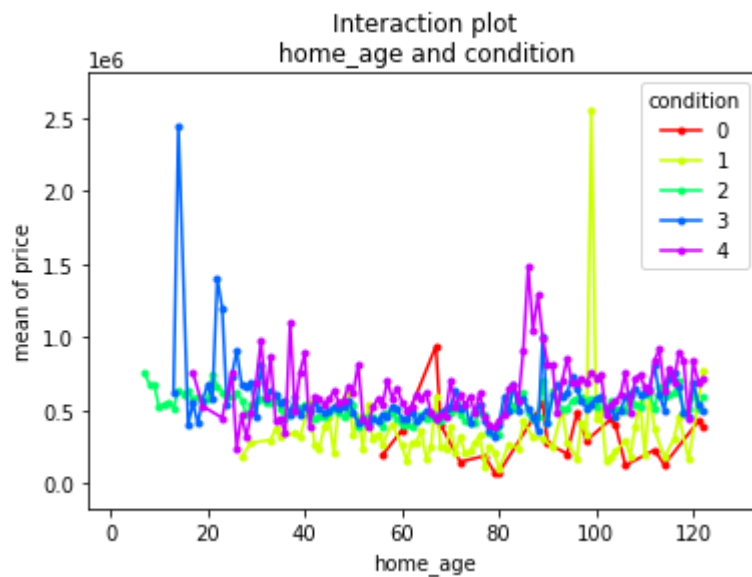
```
In [24]: import pandas as pd
from statsmodels.graphics.factorplots import interaction_plot

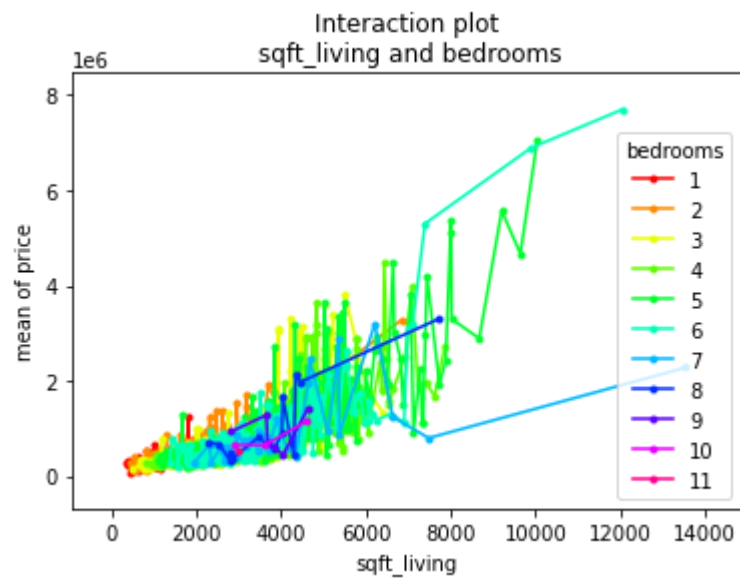
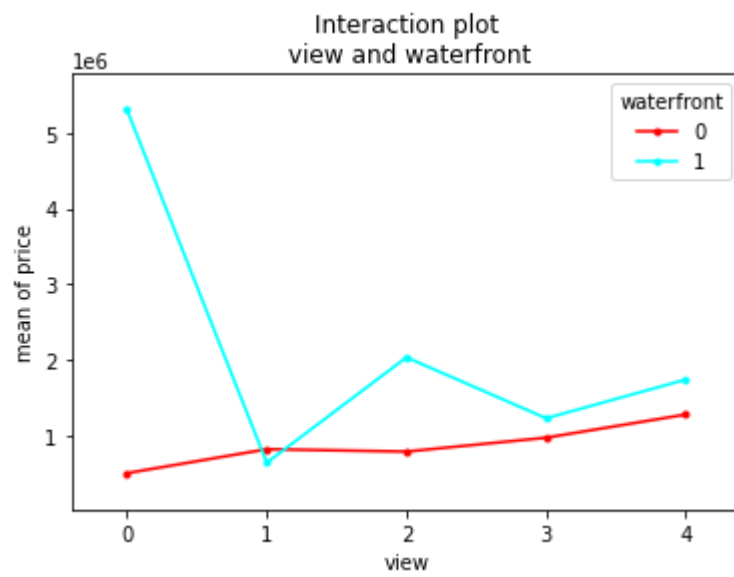

# Interaction of home age and condition
interaction_analysis(df_clean, 'home_age', 'condition')

# Interaction of bathrooms and bedrooms
interaction_analysis(df_clean, 'bathrooms', 'bedrooms')

# interaction of view and waterfront
interaction_analysis(df_clean, 'view', 'waterfront')

# Interaction of sqft_living and bedrooms
interaction_analysis(df_clean, 'sqft_living', 'bedrooms')
```





## Interaction Effects Interpretation



On an interaction plot, parallel lines indicate that there is no interaction effect while different slopes and lines that cross suggest that an interaction may be present. `home_age` and `condition`, `bathrooms` and `bedrooms`, and `sqft_living` and `bedrooms` each show interaction effects. This means a third variable influences the relationship between a dependent and independent variable. The relationship changes depending on the value of the third variable.

`view` and `waterfront` shows a slight interactive effect but only for view of 1 with no waterfront property. This may be because the view variable is measuring the view of the mountains in Washington and not the view of a waterfront property. This data also may be too subjective or variable from different observers.

## Modeling and Regression Results

In this section I run the linear regression analysis and report the results.

### Residual plots

A residual plot is produced for the independent variables with the highest correlation to `price`. Residual plots show the residual error plotted against the actual sale price. This will allow me to assess heteroskedasticity as residual plots will have a random pattern around 0 with homoskedasticity but a cone shape pattern with heteroskedasticity.

```
In [25]: # Residual plot
def residual_plot(data, X, Y, xlogged=False, ylogged=False):
    ...
    Residual Plot Function
    Input:
        data = Pandas dataframe
        X: independent variable
        Y: dependent variable
        logged: Is the axis logged?
    Output:
        Residual plot
    ...

    x = data[X]
    y = data[Y]

    # set label for x-axis if ind var is Logged
    if xlogged:
        label_x = f'log({X})'
    else: # not Logged
        label_x = f'{X}'
    # set label for y-axis if dep var is Logged
    if ylogged:
        label_y = f'log({Y})\n$USD'
    else: # not Logged
        label_y = f'{Y}\n$USD'

    # make plots
```

```

fig, (ax1, ax2) = plt.subplots(1,2, figsize=(15,4))
ax1.scatter(x, y)
m, b = np.polyfit(x, y, 1) # regression line
ax1.plot(x, m*x+b, color='red') # plot regression line
ax1.set(xlabel=label_x, ylabel=label_y)
sns.residplot(x=x, y = y, ax=ax2) # residual plot
# title of scatterplot
ax1.title.set_text(f'Scatterplot: {X} versus {Y}')
# title of residual plot
ax2.title.set_text(f'Residual plot: {X} versus {Y}')
ax2.set(ylabel='residual', xlabel=label_x) # residual plot y-axis label

```

## Linear Regression Helper Function

```

In [26]: def lin_reg_model(data, features, model_name, formula):
    ...
    Runs OLS linear regression
    Input:
        - data: clean data
        - features: independent variables that will be included in model
        - formula: regression formula in R-style
        - model_name: Name you will call the model (ex. Model 1, Model 2)
    Output:
        - OLS summary
        - Residual QQ-plot
        - OLS model object
        - Prediction interval
        - R-squared
        - Root Mean Squared Error
    ...

    target = 'price'
    y = data[target] # outcome data
    X = data[features]

    # Linear Regression using statsmodel library
    data = sm.add_constant(data)
    model = sm.OLS.from_formula(formula=formula, data=data).fit()

    # Predict values from the model
    y_predict = model.predict(X)

    # Create K-Fold cross-validation object
    kf = KFold(n_splits=5, shuffle=True) #K-Fold of 5, shuffle data, 20% test data

    # Regression model using sklearn
    lm = LinearRegression()

    # Cross-validated R-squared calculation
    r2_scores = cross_val_score(lm, X, y, cv=kf, scoring = 'r2')
    r2 = np.mean(np.absolute(r2_scores)) # calculate the mean r-squared

    # Calculate prediction interval
    sum_errs = np.sum((y - y_predict) ** 2) # Sum of errors
    stdev = np.sqrt(1/(len(y)-2) * sum_errs) # Standard deviation
    interval = 1.96 * stdev # Prediction interval

    # Cross-validated Root Mean Squared Error

```

```

scores = cross_val_score(lm, X, y, cv=kf, scoring = 'neg_mean_squared_error')
RMSE = np.sqrt(np.mean(np.absolute(scores)))

# Plotting
fig, ax = plt.subplots(1, 2, figsize=(15,4))

# Residual plot
sns.regplot(x=model.fittedvalues, y=model.resid, ax=ax[0], line_kws={'color':'r'})
ax[0].title.set_text('Residual plot of fitted values')
ax[0].set(ylabel='residuals', xlabel='fitted values')
ax[1].title.set_text('Residual QQ-plot')

# Plot residual qq-plot
sm.graphics.qqplot(model.resid, dist=stats.norm, line='45', fit=True, ax=ax[1])
plt.show() # see https://github.com/statsmodels/statsmodels/issues/5493 for bug
#fig.suptitle(f'{model_name} Residual QQ plot')

print(model.summary())
print('\n')
print(f'R-squared: {r2:.2f}')
print(f'Prediction Interval: {interval:.2f}')
print(f'Root Mean Squared Error: {RMSE:.2f}')

return model, r2, interval, RMSE

```


## M1 Results

```

In [27]: # Model 1 features (highest correlated)
model_1_features = corr_with_price['Features'].values
formula_1 = 'price ~' + '+' .join(model_1_features) # model formula

# get regression model results, prediction interval, and RMSE score
model_1_reg, model_1_r2, model_1_pi, m1_RMSE = lin_reg_model(
    df_clean, model_1_features, 'Model 1', formula_1)

```



```

                                OLS Regression Results
=====
Dep. Variable:                  price    R-squared:                  0.541
Model:                          OLS      Adj. R-squared:             0.541
Method:                        Least Squares  F-statistic:                8494.
Date:                          Tue, 09 Aug 2022  Prob (F-statistic):      0.00
Time:                          02:00:50    Log-Likelihood:             -2.9897e+05
No. Observations:                21597    AIC:                       5.980e+05
Df Residuals:                    21593    BIC:                       5.980e+05
Df Model:                        3
Covariance Type:                nonrobust
=====
               coef      std err          t      P>|t|      [0.025      0.975]
-----
Intercept    -6.564e+05    1.36e+04    -48.298    0.000    -6.83e+05    -6.3e+05
sqft_living    234.5900      4.039      58.075    0.000     226.672     242.508
grade         1.108e+05    2325.608     47.637    0.000     1.06e+05     1.15e+05
sqft_above    -78.0959      4.427     -17.642    0.000     -86.773     -69.419
=====
Omnibus:                17102.886    Durbin-Watson:              1.976
Prob(Omnibus):           0.000    Jarque-Bera (JB):           1062513.676
Skew:                    3.332    Prob(JB):                   0.00
Kurtosis:                36.709    Cond. No.                   2.43e+04
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.43e+04. This might indicate that there are strong multicollinearity or other numerical problems.

R-squared: 0.54

Prediction Interval: 487671.26

Root Mean Squared Error: 248906.83

## Interpretation of M1 price variable

The p-values indicate that each of the variables chosen has a statistical significant relationship with the dependent variable, the  $R^2$  value is low at 0.54.

The residual plot shows a cone like pattern indicating the homoskedasticity assumption is not valid. Also the assumption that the error term is normally distributed is not met as the residuals are not normally distributed. To make accurate inferences these assumptions need to be met.

As the dependent variable, `price` is not normally distributed, I will log transform `price` to keep the variance constant and not change proportionally with the magnitude of `price`.

## Log transforming `price` variable

The distribution of the dependent variable, `price`, is normally distributed after being log transformed.

```
In [28]: # Log transform price data
```



```
df_clean['price'] = np.log(df_clean['price'])
```

```
In [29]: # Histogram of log transformed price independent variable  
hist_plot(df_clean, 'price')
```


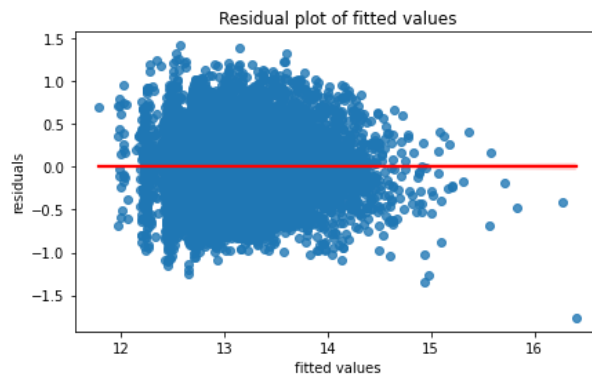
EDA of price variable

Skewness: 0.4310041773299232

Kurtosis: 0.691048515911131



```
In [30]: # Regression analysis using log transformed dependent variable  
# get regression model results, prediction interval, and RMSE score  
model_1_reg, model_1_r2, model_1_pi, m1_RMSE = lin_reg_model(  
    df_clean, model_1_features, 'Model 1', formula_1)
```



```

                                OLS Regression Results
=====
Dep. Variable:                  price    R-squared:                  0.564
Model:                          OLS      Adj. R-squared:             0.564
Method:                        Least Squares  F-statistic:                9322.
Date:                          Tue, 09 Aug 2022  Prob (F-statistic):      0.00
Time:                          02:00:52    Log-Likelihood:             -7820.6
No. Observations:                21597    AIC:                       1.565e+04
Df Residuals:                    21593    BIC:                       1.568e+04
Df Model:                        3
Covariance Type:                nonrobust
=====
                                coef      std err          t      P>|t|      [0.025      0.975]
-----
Intercept                11.0806         0.019    583.615     0.000     11.043     11.118
sqft_living              0.0003      5.64e-06    53.520     0.000     0.000     0.000
grade                   0.2055         0.003     63.266     0.000     0.199     0.212
sqft_above             -0.0001      6.18e-06   -21.209     0.000    -0.000    -0.000
=====
Omnibus:                   37.044    Durbin-Watson:             1.973
Prob(Omnibus):             0.000    Jarque-Bera (JB):          37.179
Skew:                      0.101    Prob(JB):                  8.45e-09
Kurtosis:                  3.029    Cond. No.:                 2.43e+04
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.43e+04. This might indicate that there are strong multicollinearity or other numerical problems.

R-squared: 0.56

Prediction Interval: 0.68

Root Mean Squared Error: 0.35

## M1 Interpretation

The p-values ( $\alpha < 0.05$ ) indicate that each of the variables chosen has a statistical significant relationship with the dependent variable. The residual plot has a random pattern and the QQ-plot of the residuals indicate that they are normally distributed. Log transforming `price` has worked and the assumptions are met. There does seem to be a high level of multicollinearity.

The linear function of M1 is:

$$\hat{y}_{sales} = 0.0003X_{sqft_{living}} + 0.201X_{grade} - 0.0001X_{sqft_{above}} + 11.08$$

To interpret the coefficients we need to take into account that `price` has been log transformed (3).

$$(e^{\beta_0} - 1) \times 100, \text{ where } \beta_0 \text{ is the coefficient}$$

The coefficient for `grade` of a home is 0.201. This means for every one-unit increase in the `grade` of the home, there is a 22.8% increase in sales price. The coefficient for `sqft_living` is 0.0003. This means for every square foot increase in the living space in a house there is a



0.03% increase in sale price. Lastly, the coefficient for `sqft_above` is 0.0001 so for every square foot increase in the above ground area of a house there is a 0.01% reduction in sale price. Interestingly, the area of the space above the home is penalized in this model. This may be due to multicollinearity between the `sqft_living` and `sqft_above`. These values are so low as to be insignificant.

$R^2$  value is low at 0.56, RMSE is 0.35, and prediction interval is 0.68.

## Model 2 Results

```
In [31]: # Model 2 features (Stepwise backward Design)
model_2_features = backward_elimination(df_clean, df_clean[['price']])
formula_2 = 'price ~' + '+' .join(model_2_features) # model formula

# get regression model results, prediction interval, and RMSE score
model_2_reg, model_2_r2, model_2_pi, m2_RMSE = lin_reg_model(df_clean, model_2_features)
```



# OLS Regression Results

```

=====
Dep. Variable:          price      R-squared:                0.877
Model:                  OLS        Adj. R-squared:           0.876
Method:                 Least Squares  F-statistic:            1938.
Date:                  Tue, 09 Aug 2022  Prob (F-statistic):      0.00
Time:                  02:00:55      Log-Likelihood:         5818.9
No. Observations:      21597        AIC:                   -1.148e+04
Df Residuals:          21517        BIC:                   -1.084e+04
Df Model:              79
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	11.0932	0.016	692.273	0.000	11.062	11.125
bathrooms	0.0383	0.003	12.792	0.000	0.032	0.044
sqft_living	9.116e-05	5.44e-06	16.751	0.000	8.05e-05	0.000
sqft_lot	6.304e-07	3.32e-08	18.998	0.000	5.65e-07	6.95e-07
floors	-0.0293	0.004	-8.027	0.000	-0.036	-0.022
waterfront	0.4626	0.017	27.368	0.000	0.429	0.496
view	0.0610	0.002	30.593	0.000	0.057	0.065
condition	0.0587	0.002	26.477	0.000	0.054	0.063
grade	0.0921	0.002	44.088	0.000	0.088	0.096
sqft_above	0.0001	6.12e-06	18.722	0.000	0.000	0.000
basement	0.0437	0.005	9.260	0.000	0.034	0.053
sqft_living15	8.428e-05	3.32e-06	25.365	0.000	7.78e-05	9.08e-05
home_age	0.0018	0.000	14.790	0.000	0.002	0.002
yr_since_reno	-0.0013	0.000	-11.027	0.000	-0.002	-0.001
zipcode_98002	-0.0494	0.014	-3.471	0.001	-0.077	-0.022
zipcode_98004	1.0693	0.012	89.714	0.000	1.046	1.093
zipcode_98005	0.7004	0.015	45.408	0.000	0.670	0.731
zipcode_98006	0.5915	0.010	58.200	0.000	0.572	0.611
zipcode_98007	0.6236	0.017	37.693	0.000	0.591	0.656
zipcode_98008	0.6205	0.012	50.433	0.000	0.596	0.645
zipcode_98010	0.2268	0.019	11.751	0.000	0.189	0.265
zipcode_98011	0.4266	0.014	29.751	0.000	0.398	0.455
zipcode_98014	0.2828	0.018	16.061	0.000	0.248	0.317
zipcode_98019	0.3132	0.014	21.610	0.000	0.285	0.342
zipcode_98023	-0.0547	0.010	-5.538	0.000	-0.074	-0.035
zipcode_98024	0.3948	0.022	18.346	0.000	0.353	0.437
zipcode_98027	0.4779	0.011	44.852	0.000	0.457	0.499
zipcode_98028	0.3997	0.012	32.598	0.000	0.376	0.424
zipcode_98029	0.5740	0.012	48.820	0.000	0.551	0.597
zipcode_98030	0.0367	0.013	2.882	0.004	0.012	0.062
zipcode_98031	0.0540	0.012	4.350	0.000	0.030	0.078
zipcode_98032	-0.0552	0.017	-3.165	0.002	-0.089	-0.021
zipcode_98033	0.7527	0.010	72.111	0.000	0.732	0.773
zipcode_98034	0.5155	0.010	53.736	0.000	0.497	0.534
zipcode_98038	0.1547	0.009	16.521	0.000	0.136	0.173
zipcode_98039	1.1889	0.027	44.069	0.000	1.136	1.242
zipcode_98040	0.8255	0.013	65.594	0.000	0.801	0.850
zipcode_98042	0.0402	0.010	4.207	0.000	0.021	0.059
zipcode_98045	0.3099	0.014	22.810	0.000	0.283	0.337
zipcode_98052	0.6092	0.009	64.215	0.000	0.591	0.628
zipcode_98053	0.5565	0.011	51.626	0.000	0.535	0.578
zipcode_98055	0.1139	0.013	9.087	0.000	0.089	0.139
zipcode_98056	0.2907	0.011	27.302	0.000	0.270	0.312
zipcode_98058	0.1405	0.010	13.787	0.000	0.121	0.160
zipcode_98059	0.3103	0.010	30.544	0.000	0.290	0.330
zipcode_98065	0.3713	0.012	31.010	0.000	0.348	0.395



zipcode_98070	0.3129	0.018	17.111	0.000	0.277	0.349
zipcode_98072	0.4663	0.012	37.410	0.000	0.442	0.491
zipcode_98074	0.5258	0.010	50.175	0.000	0.505	0.546
zipcode_98075	0.5187	0.011	45.402	0.000	0.496	0.541
zipcode_98077	0.4149	0.014	28.818	0.000	0.387	0.443
zipcode_98102	0.8992	0.019	46.257	0.000	0.861	0.937
zipcode_98103	0.7773	0.010	78.771	0.000	0.758	0.797
zipcode_98105	0.8937	0.014	64.466	0.000	0.867	0.921
zipcode_98106	0.2966	0.012	25.509	0.000	0.274	0.319
zipcode_98107	0.7925	0.013	61.047	0.000	0.767	0.818
zipcode_98108	0.3206	0.015	21.721	0.000	0.292	0.350
zipcode_98109	0.9393	0.019	49.576	0.000	0.902	0.976
zipcode_98112	0.9905	0.013	75.090	0.000	0.965	1.016
zipcode_98115	0.7785	0.010	79.822	0.000	0.759	0.798
zipcode_98116	0.7169	0.012	60.352	0.000	0.694	0.740
zipcode_98117	0.7658	0.010	76.917	0.000	0.746	0.785
zipcode_98118	0.4185	0.010	41.460	0.000	0.399	0.438
zipcode_98119	0.9274	0.015	60.927	0.000	0.898	0.957
zipcode_98122	0.7558	0.013	59.629	0.000	0.731	0.781
zipcode_98125	0.5371	0.011	49.958	0.000	0.516	0.558
zipcode_98126	0.4999	0.011	43.581	0.000	0.477	0.522
zipcode_98133	0.4257	0.010	42.415	0.000	0.406	0.445
zipcode_98136	0.6385	0.013	49.597	0.000	0.613	0.664
zipcode_98144	0.6166	0.012	52.688	0.000	0.594	0.640
zipcode_98146	0.2485	0.012	20.230	0.000	0.224	0.273
zipcode_98148	0.1315	0.025	5.227	0.000	0.082	0.181
zipcode_98155	0.4002	0.010	38.690	0.000	0.380	0.421
zipcode_98166	0.2795	0.013	21.680	0.000	0.254	0.305
zipcode_98168	0.0438	0.013	3.465	0.001	0.019	0.069
zipcode_98177	0.5603	0.013	43.356	0.000	0.535	0.586
zipcode_98178	0.1148	0.013	8.982	0.000	0.090	0.140
zipcode_98188	0.0745	0.017	4.433	0.000	0.042	0.107
zipcode_98198	0.0369	0.012	2.987	0.003	0.013	0.061
zipcode_98199	0.8144	0.012	67.498	0.000	0.791	0.838

```
=====
Omnibus:                1934.611    Durbin-Watson:           2.010
Prob(Omnibus):          0.000    Jarque-Bera (JB):       8714.901
Skew:                   -0.339    Prob(JB):               0.00
Kurtosis:                6.037    Cond. No.               1.62e+06
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.62e+06. This might indicate that there are strong multicollinearity or other numerical problems.

R-squared: 0.88

Prediction Interval: 0.36

Root Mean Squared Error: 0.19

## M2 Interpretation

There are many more independent variables used in M2 as compared to M1, which may overfit the data. This model used all the independent variables except for some of the Zip codes after recoding Zip code to dummy variables. `price` is still log transformed. Though there are many variables, the p-values indicate that each of the variables chosen has a statistical significant


relationship with the dependent variable, `price`. The residual plot has a random pattern so the homoskedasticity assumption is met. The QQ-plot of the residuals indicate that they are not normally distributed.  $R^2$  is 0.88 for M2. This is tremendous improvement over M1. RMSE is 0.19 and the prediction interval is 0.36.

The coefficient for `condition` for M2 is 0.0587. This means that for ever one-unit increase in condition of a house the sale price is expected to increase on average 6%. Adding an additional full bathroom would increase the sale price of a home by about 3.9%.

M2 can tell us which customers to target when marketing. Waterfront properties have a sale price over 59% more homes not on water. A 1-unit increase in the score of a view increases the sale price by about 6%. Some of the highest coefficients with M2 is in the Zip code variables. For example, with all other independent variables held constant a home in 98039 would sell for 228% more than Zip code 98003. The real estate agents should avoid marketing to property owners in 98002, 98003, 98023, 98032, 98042, and 98198 as these sell for much less than other Zip codes in King County, WA as these sell for much less than other Zip codes in King County, WA.

## M3 Results

```
In [32]: model_1_features_del_sqftabove = np.delete(model_1_features, 2)
model_3_features = np.append(model_1_features_del_sqftabove, ['bedrooms', 'home_age',
formula_3 = 'price ~' + '+'.join(model_3_features) + '+ home_age*condition + sqft_livi
model_3_reg, model_3_r2, model_3_pi, m3_RMSE = lin_reg_model(df_clean, model_3_feature
```



```

                                OLS Regression Results
=====
Dep. Variable:                  price    R-squared:                  0.625
Model:                          OLS      Adj. R-squared:             0.625
Method:                        Least Squares  F-statistic:                5141.
Date:                          Tue, 09 Aug 2022  Prob (F-statistic):      0.00
Time:                          02:00:56    Log-Likelihood:            -6199.3
No. Observations:              21597      AIC:                      1.241e+04
Df Residuals:                  21589      BIC:                      1.248e+04
Df Model:                      7
Covariance Type:               nonrobust
=====
===
                                coef      std err          t      P>|t|      [0.025      0.9
75]
-----
---
Intercept                    10.6119      0.035      306.963      0.000      10.544      10.
680
sqft_living                   0.0003     9.76e-06      28.011      0.000      0.000      0.
000
grade                         0.2421      0.003      77.287      0.000      0.236      0.
248
bedrooms                     0.0064      0.005       1.193      0.233     -0.004      0.
017
home_age                     0.0011      0.000       3.344      0.001      0.000      0.
002
condition                    -0.0718      0.009     -7.563      0.000     -0.090     -0.
053
home_age:condition           0.0016      0.000      12.406      0.000      0.001      0.
002
sqft_living:bedrooms -1.339e-05     2.19e-06     -6.115      0.000     -1.77e-05     -9.1e
-06
=====
Omnibus:                     50.969    Durbin-Watson:              1.961
Prob(Omnibus):                0.000    Jarque-Bera (JB):          56.642
Skew:                         -0.077    Prob(JB):                  5.02e-13
Kurtosis:                     3.197    Cond. No.                  1.48e+05
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.48e+05. This might indicate that there are strong multicollinearity or other numerical problems.

R-squared: 0.62

Prediction Interval: 0.63

Root Mean Squared Error: 0.32

## M3 Interpretation

M3 includes interaction effects along with main effects. The interaction effect of **bathrooms** and **bedrooms** was not included as the p-value was above 0.05, indicating that the interaction with **price** is not statistically significant. Likewise, the variable **bedrooms** is not have a statistically significant main interaction with **price**. The residual plot has a random pattern so

there is homoskedasticity. The QQ-plot of the residuals indicate that they are normally distributed. These results show the model meets the assumptions of OLS. There is multicollinearity between some of the variables.

$R^2$  is 0.62. This is an improvement over M1 but is less than M2. RMSE is 0.32 and the prediction interval is 0.63.

The coefficient for `sqft_living` is 0.0003, the same as M1. The coefficient for `condition` with M3 is -0.077 so for each one-unit increase in the condition of a house the sale price is expected to **decrease** about 8% on average.

## Conclusions

### Model Analysis and Comparisons

I will use coefficient of determination ( $R^2$ ), prediction intervals, and root means squared error (RMSE) to compare the models and determine which model meets the needs of our client. This includes the most accurate predictive power with a tight range in possibilities.

```
In [33]: # Table of model's R-squared, PI, and Cross-validated RMSE
r2 = [model_1_reg.rsquared, model_2_reg.rsquared, model_3_reg.rsquared]
pi = [model_1_pi, model_2_pi, model_3_pi]
rmse = [m1_RMSE, m2_RMSE, m3_RMSE]
results = pd.DataFrame({'R_squared': r2,
                        'PI': pi,
                        'RMSE': rmse},
                        index = ['M1', 'M2', 'M3'])
print('Table 5: Regression Results Table')
results
```

Table 5: Regression Results Table

```
Out[33]:
```

	R_squared	PI	RMSE
<b>M1</b>	0.564301	0.681244	0.347598
<b>M2</b>	0.876794	0.362265	0.185685
<b>M3</b>	0.625043	0.631976	0.323954

### Coefficient of Determination ( $R^2$ )

$R^2$  is a statistical estimate of how close the observed data is to the regression line of each model. It is the proportion of variation in the dependent variable that is predictive from the independent variable. I measured  $R^2$  using 5-fold cross-validation.

$$R^2 = 1 - \frac{\text{Residual Sum of Squares (RSS)}}{\text{Total Sum of Squares (TSS)}}$$

$$= 1 - \frac{\sum (y_i - \hat{y})^2}{\sum (y_i - \bar{y})^2}$$

As Table 5 above shows, the model with the highest  $R^2$  is M2 at 0.88. This is a great score for a predictive model, the higher the better. M1 has the lowest  $R^2$  score at 0.55 and M3 is 0.63.

## Prediction Interval

A prediction interval (PI) is the range where a single new observation is likely to fall given specific values of the independent variables. The prediction interval can be used to assess if the predictions are sufficiently in a narrow range to satisfy the client's requirement. Prediction intervals can be compared across models. Smaller intervals indicate tighter predictive range. Large prediction intervals tell us the model could have a wide range in its predictions and would not meet the client's needs.

The prediction interval is calculated by,

$$PI = 1.96 \times s,$$

where  $s$  is the sample standard deviation calculated by

$$s = \sqrt{\frac{1}{N-2} \times RSS}$$

M2 also has the lowest predictive intervals. A prediction interval is the range where a single new observation is likely (95%) to fall given specific values of the independent variables. The smaller the predictive interval the more confidence the true sale price is in that region. M1 and M2 had reasonably similar prediction intervals and twice the value as M2.

## Root Mean Squared Error

RMSE is a measure of the mean error rate of a regression model that penalizes larger errors. The smaller the RMSE value, the closer the fitted line from the linear equation is to the actual data. It is the square root of the average squared difference between the predicted dependent value and the actual values in the dataset. Like Mean Squared Error (MSE), this statistic squares the residual error before it is averaged, which gives a high weight to large errors, but because the square root is taken, the statistic is in the same units as the dependent variable, sale price (\$USD). The lower the score of RMSE the closer the model fits the data.

$$RMSE = \sqrt{\frac{1}{n} \sum_{i=1}^n (\hat{y}_i - y_i)^2},$$

where  $(\hat{y}_1, \hat{y}_2, \dots, \hat{y}_n)$  are the predicted values,  
 $(y_1, y_2, \dots, y_n)$  are the observed values  
and  $(n)$  is the number of observations

RMSE is estimated using K-Fold cross validation. M2 is almost half the RMSE score of M1 and M3 indicating it produces better prediction between the actual and predicted values.

## Discussion on the best model

Although M2 does not meet all the assumptions of OLS linear regression analysis, it is the best model of the three as it has a high  $R^2$ , low RMSE, and low PI. M2 also has the most interpretable variables. M1's is almost entirely dependent on the `grade` of the house. The `grade` of a home cannot be changed unless the house is completely remodeled. This model would not help our clients as it has low predictive power and does not have variables that could be improved to increase the sale price of the house.

According to M3, increasing the condition of a house will decrease the sale price of a home, which doesn't make a lot of sense.

M2 has many variables that can be interpreted to improve the sale price of the house like improving the condition and number of bathrooms. M2 is driven heavily by the location of the house in the form of the Zip code. These coefficients will be used to make actionable recommendations to real estate agents for targeted marketing.

## Prediction Application

The prediction of Model 2 seems the most reliable of the three models. Below is the `price_predictor` function that input one of the three models, the new home features, and the prediction interval. This function returns the predicted house sale price and 95% prediction interval range according to that model and new data.

```
In [34]: def price_predictor(model_no, new_data, pi):
    """
    This function takes in a model and Pandas series object and returns
    an estimated price range for the house to be listed.

    Model 1: price ~ ln(sqft_living) + ln(sqft_above) + grade

    Model 2: price ~ bedrooms+bathrooms+sqft_living+sqft_lot+floors+waterfront+
              view+condition+grade+sqft_above+basement+sqft_living15+sqft_lot15
              home_age+yr_since_reno+zipcode

    Model 3: price ~ sqft_living+grade+sqft_above+ home_age*condition + sqft_living*be

    Input:
        model_no: Statsmodel linear regression model results
        new_data: Pandas Series with variables needed for the regression model
        pi: prediction interval for that model

    Output:
        Predicted sale price
        Predicted sale price range
    """
    price_ln = model_no.predict(new_data).values[0] # predicted price
    price = round(np.exp(price_ln))
    price_low = round(np.exp(price_ln - pi)) # prediction lower bound
    price_up = round(np.exp(price_ln + pi)) # prediction upper bound

    return f'Predicted price: {price}, range: {price_low} - {price_up}'
```

## Demonstration

```
In [35]: #data = df_clean.iloc[8677] # for testing purposes
```

```
In [48]: data = df_clean.sample().squeeze() # Randomly sample data for prediction  
data.head(10) # new house features
```

```
Out[48]: price          13.122363  
bedrooms          3.000000  
bathrooms          2.500000  
sqft_living    1210.000000  
sqft_lot       1200.000000  
floors           3.000000  
waterfront       0.000000  
view             0.000000  
condition        2.000000  
grade            8.000000  
Name: 20464, dtype: float64
```

```
In [49]: # Model 1 prediction  
data_new = data.drop('price') # Remove actual price  
print(price_predictor(model_1_reg, data_new, model_1_pi)) # Run prediction function  
sale_price = round(np.exp(data['price'])) # assign actual sale price  
print(f'Actual sale price: {sale_price}') # print out actual sale price
```

Predicted price: 413149, range: 209048 - 816520  
Actual sale price: 500000

```
In [50]: # Model 2 Prediction function  
print(price_predictor(model_2_reg, data_new, model_2_pi))  
print(f'Actual sale price: {sale_price}')
```

Predicted price: 484124, range: 336998 - 695483  
Actual sale price: 500000

```
In [51]: # Model 3 prediction function  
print(price_predictor(model_3_reg, data_new, model_3_pi))  
print(f'Actual sale price: {sale_price}')
```

Predicted price: 350551, range: 186332 - 659499  
Actual sale price: 500000

## Recommendations

### Summary

- Our client wants to be able to predict sales price, identify where to market in King County, WA, and how customers can improve their home to increase sale price.
- Ordinary least squares linear regression was used to create three models.
- The three models were compared using  $R^2$ , Prediction Intervals (PI), and Root Mean Squared Error (RMSE).
- Model 2 (M2) is the best model as it has the best predictive capabilities, R-squared of 0.88, low RMSE and PI, though the error is not normally distributed.

## Actionable Recommendations

1. Improving the condition of a house by one-unit will increase the sale price by about 6%.
2. Adding an additional full bathroom would increase the sale price of a house by about 3.9%.
3. Marketing should be focused throughout King County, WA except in Zip codes 98002, 98003, 98023, 98032, 98042, and 98198 as the value of the homes sold in these Zip codes are well under the rest of King County, WA.
4. Market real estate services toward owners of waterfront properties as these sell for 59% more than homes not waterfront.

## Next Steps

This model could be improved to make better predictions by adding more data and additional variables, such as crime rate in the geographic location of the home, the zoned public school ranking, and time the house was on the market until it was sold. The GPS coordinates of the sold house could be used to collect the first two of these variables. The Multiple Listing Service may be a source for more recent data and on how long a house was on the market from day of listing to closing date.

Another source of data could be in the internal data of our brokerage client. They possibly have data of properties they have sold or bid on, this would include the data of the asking and bidding price of the property.

## References

1. Kaggle, [Kaggle, House Sales in King County, USA](#)
2. Albert, [Key Assumptions of OLS: Econometrics Review](#)
3. University of Virginia, [Interpreting Log Transformations in a Linear Model](#)